



Unsupervised Learning: Clustering with K-Means

Supervised learning




Unsupervised learning



Training set: $\{\underline{x^{(1)}}, \underline{x^{(2)}}, x^{(3)}, \dots, \underline{x^{(m)}}\}$ \leftarrow

Unsupervised learning




Clustering algorithm

Training set: $\{\underline{x^{(1)}}, \underline{x^{(2)}}, \underline{x^{(3)}}, \dots, \underline{x^{(m)}}\}$ ←

Clustering:

Finding a way to divide a dataset into groups ('clusters')

- Data points within the same cluster should be 'close' or 'similar' in some way.
- Data points in different clusters should be 'far apart' or 'different'
- Clustering algorithms output a cluster membership index for each data point:
 - *Hard clustering: each data point belongs to exactly one cluster*
 - *Soft (or fuzzy) clustering: each data point is assigned a weight, score, or probability of membership for each cluster*



K-means Clustering

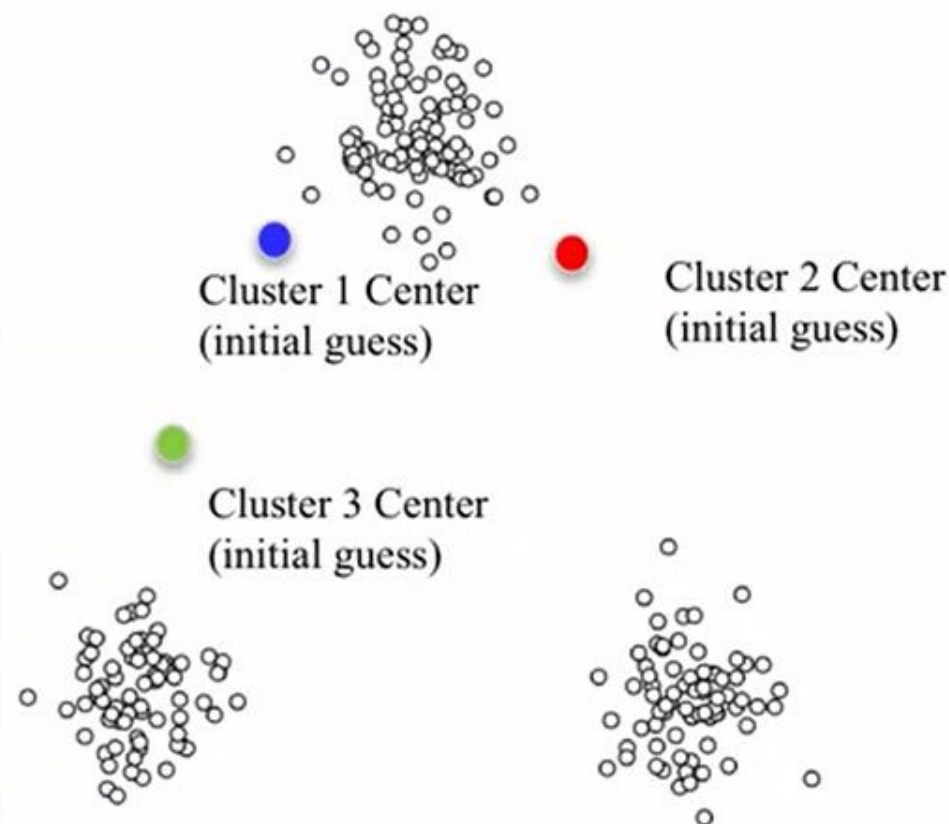
The k-means algorithm

Initialization Pick the number of clusters k you want to find.
Then pick k *random* points to serve as an initial guess for the cluster centers.


Step A Assign each data point to the nearest cluster center.


Step B Update each cluster center by replacing it with the mean of all points assigned to that cluster (in step A).


Repeat steps A and B until the centers converge to a stable solution.





Demo: <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>



























Local optima



Random initialization

For $i = 1$ to 100 {

Randomly initialize K-means.

Run K-means. Get $c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K$.

Compute cost function (distortion)

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

}


k-means Example in Scikit-Learn

```
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from adspy_shared_utilities import plot_labelled_scatter
```

```
X, y = make_blobs(random_state = 10)
```

```
kmeans = KMeans(n_clusters = 3)
kmeans.fit(X)
```

```
plot_labelled_scatter(X, kmeans.labels_, ['Cluster 1', 'Cluster 2', 'Cluster 3'])
```




Choosing K


Clustering

Unsupervised Learning


What is the right value of K?



What is the right value of K?



What is the right value of K?



Choosing the value of K

Elbow method:

