

Transform and Conquer



This group of techniques solves a problem by a *transformation* to

- ❑ a simpler/more convenient instance of the same problem (*instance simplification*)
- ❑ a different representation of the same instance (*representation change*)
- ❑ a different problem for which an algorithm is already available (*problem reduction*)

Instance simplification - Presorting



Solve a problem's instance by transforming it into another simpler/easier instance of the same problem

Presorting

Many problems involving lists are easier when list is sorted, e.g.

- ❑ **searching**
- ❑ **computing the median (selection problem)**
- ❑ **checking if all elements are distinct (element uniqueness)**

Also:

- ❑ **Topological sorting helps solving some problems for dags.**
- ❑ **Presorting is used in many geometric algorithms.**

How fast can we sort ?



Efficiency of algorithms involving sorting depends on efficiency of sorting.

Theorem (see Sec. 11.2): $\lceil \log_2 n! \rceil \approx n \log_2 n$ comparisons are necessary in the worst case to sort a list of size n by any comparison-based algorithm.

Note: About $n \log_2 n$ comparisons are also sufficient to sort array of size n (by mergesort).

Searching with presorting



Problem: Search for a given K in $A[0..n-1]$

Presorting-based algorithm:

Stage 1 Sort the array by an efficient sorting algorithm

Stage 2 Apply binary search

Efficiency: $\Theta(n \log n) + O(\log n) = \Theta(n \log n)$

Good or bad?

**Why do we have our dictionaries, telephone directories, etc.
sorted?**

Element Uniqueness with presorting



❑ Presorting-based algorithm

Stage 1: sort by efficient sorting algorithm (e.g. mergesort)

Stage 2: scan array to check pairs of adjacent elements

Efficiency: $\Theta(n \log n) + O(n) = \Theta(n \log n)$

❑ Brute force algorithm

Compare all pairs of elements

Efficiency: $O(n^2)$

❑ Another algorithm? Hashing

Searching Problem



Problem: Given a (multi)set S of keys and a search key K , find an occurrence of K in S , if any

❑ Searching must be considered in the context of:

- file size (internal vs. external)
- dynamics of data (static vs. dynamic)

❑ Dictionary operations (dynamic data):

- find (search)
- insert
- delete

Taxonomy of Searching Algorithms



❑ List searching

- **sequential search**
- **binary search**
- **interpolation search**

❑ Tree searching

- **binary search tree**
- **binary balanced trees: AVL trees, red-black trees**
- **multiway balanced trees: 2-3 trees, 2-3-4 trees, B trees**

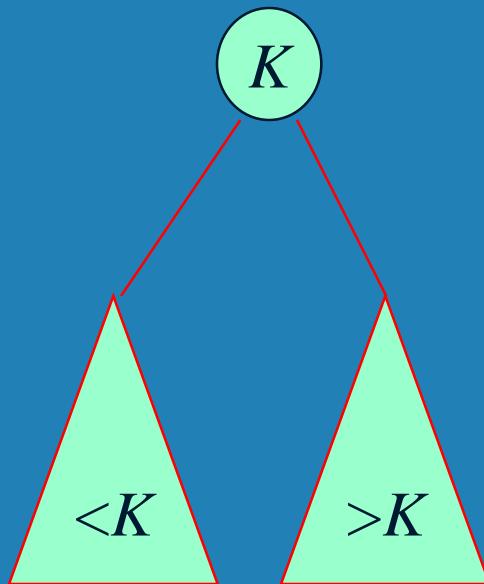
❑ Hashing

- **open hashing (separate chaining)**
- **closed hashing (open addressing)**

Binary Search Tree



Arrange keys in a binary tree with the *binary search tree property*:



Example: 5, 3, 1, 10, 12, 7, 9

Dictionary Operations on Binary Search Trees



Searching – straightforward

Insertion – search for key, insert at leaf where search terminated

Deletion – 3 cases:

deleting key at a leaf

deleting key at node with single child

deleting key at node with two children

**Efficiency depends of the tree's height: $\lfloor \log_2 n \rfloor \leq h \leq n-1$,
with height average (random files) be about $3\log_2 n$**

Thus all three operations have

- **worst case efficiency: $\Theta(n)$**
- **average case efficiency: $\Theta(\log n)$**

Bonus: inorder traversal produces sorted list

Balanced Search Trees



Attractiveness of *binary search tree* is marred by the bad (linear) worst-case efficiency. Two ideas to overcome it are:

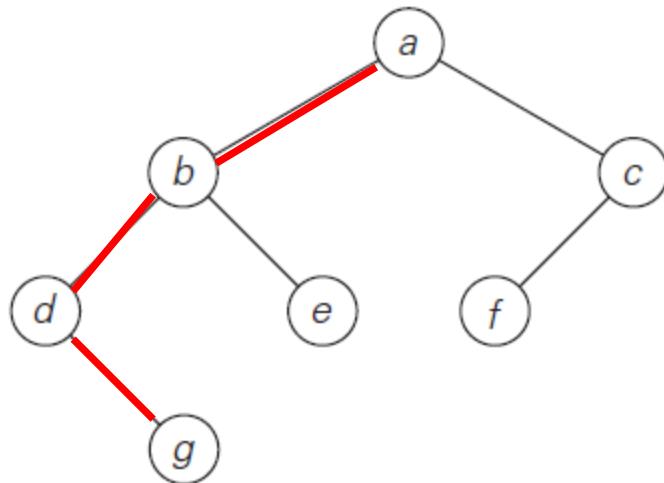
- ❑ to rebalance binary search tree when a new insertion makes the tree “too unbalanced”
 - *AVL trees*
 - *red-black trees*

- ❑ to allow more than one key per node of a search tree
 - *2-3 trees*
 - *2-3-4 trees*
 - *B-trees*

Balanced trees: AVL trees



Definition An *AVL tree* is a binary search tree in which, for every node, the difference between the heights of its left and right subtrees, called the *balance factor*, is at most 1 (with the height of an empty tree defined as -1)



$$H(a) = 3 \text{ (number of edges in the longest path)}$$

ALGORITHM *Height(T)*

```
//Computes recursively the height of a binary tree
//Input: A binary tree  $T$ 
//Output: The height of  $T$ 
if  $T = \emptyset$  return -1
else return  $\max\{\text{Height}(\text{Node.left}), \text{Height}(\text{Node.right})\} + 1$ 
```

$$H(\emptyset) = -1$$

$$H(\text{Single Node}) = 0$$

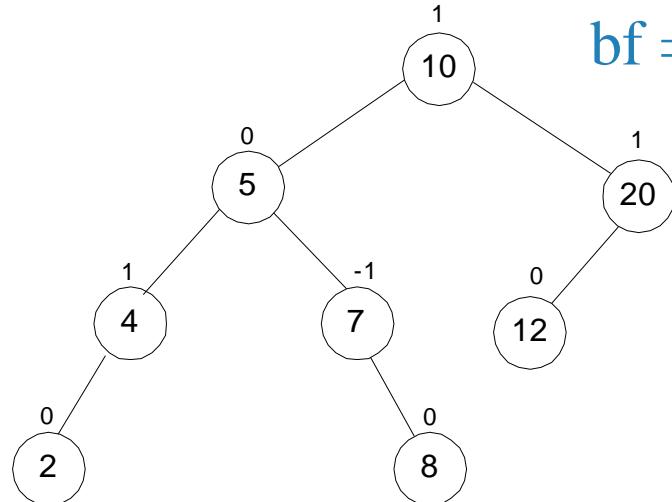
$$H(\text{Node}) = \max\{H(\text{Node.left}), H(\text{Node.right})\} + 1$$

Balanced trees: AVL trees

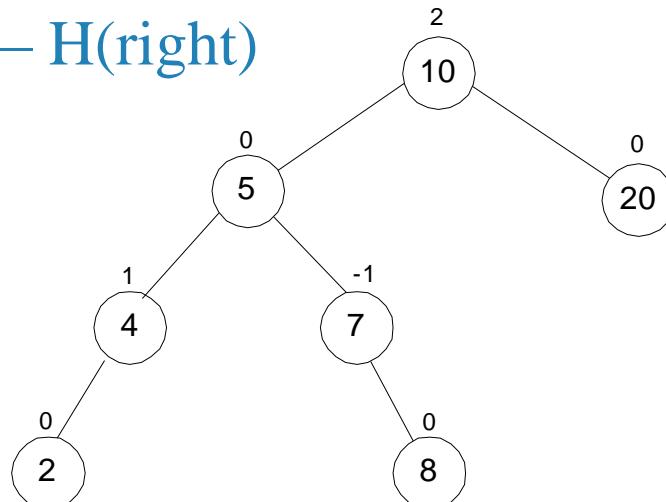


Definition An *AVL tree* is a binary search tree in which, for every node, the difference between the heights of its left and right subtrees, called the *balance factor*, is at most 1 (with the height of an empty tree defined as -1)

$$bf = H(\text{left}) - H(\text{right})$$



(a)



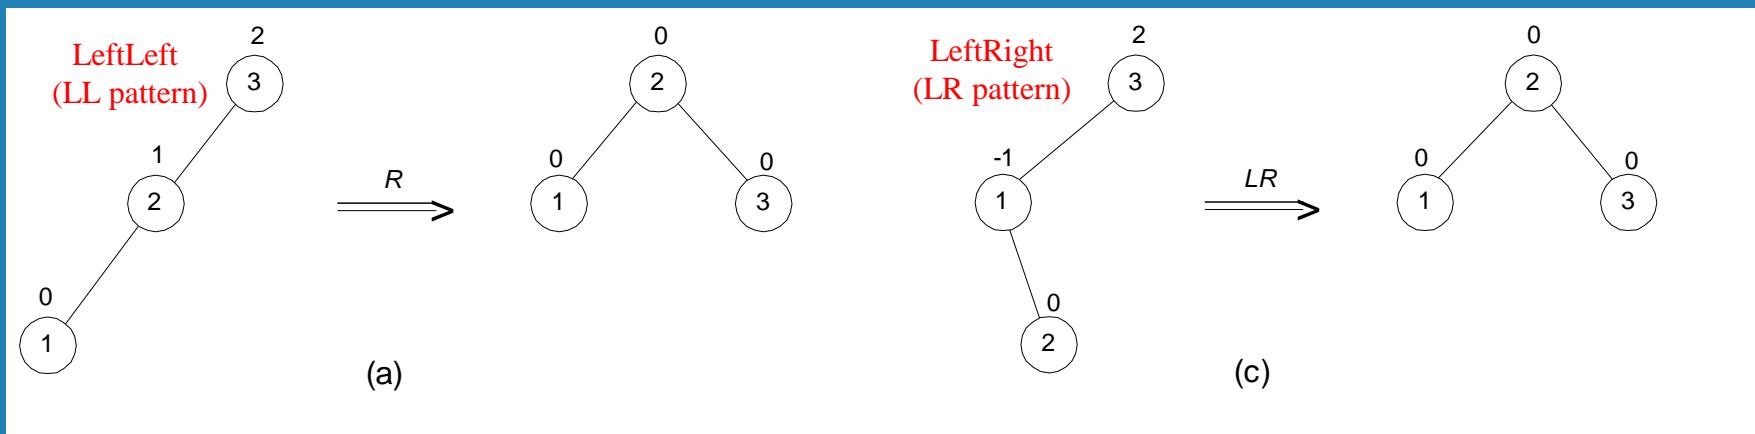
(b)

Tree (a) is an AVL tree; tree (b) is not an AVL tree

Rotations



If a key insertion violates the balance requirement at some node, the subtree rooted at that node is transformed via one of the four *rotations*. (The rotation is always performed for a subtree rooted at an “unbalanced” node closest to the new leaf.)



Single *R*-rotation

Double *LR*-rotation

Rotations

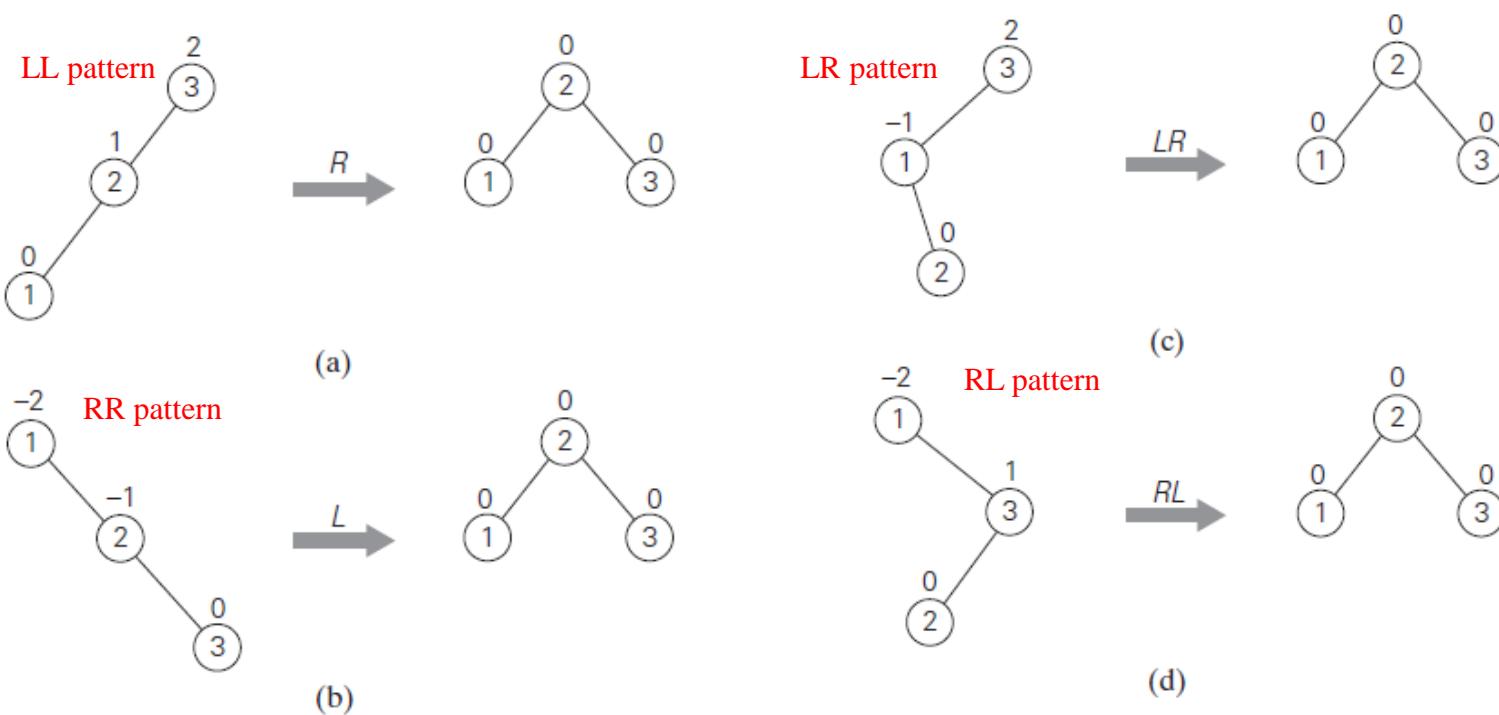


FIGURE 6.3 Four rotation types for AVL trees with three nodes. (a) Single R -rotation.
(b) Single L -rotation. (c) Double LR -rotation. (d) Double RL -rotation.

General case: Single R-rotation

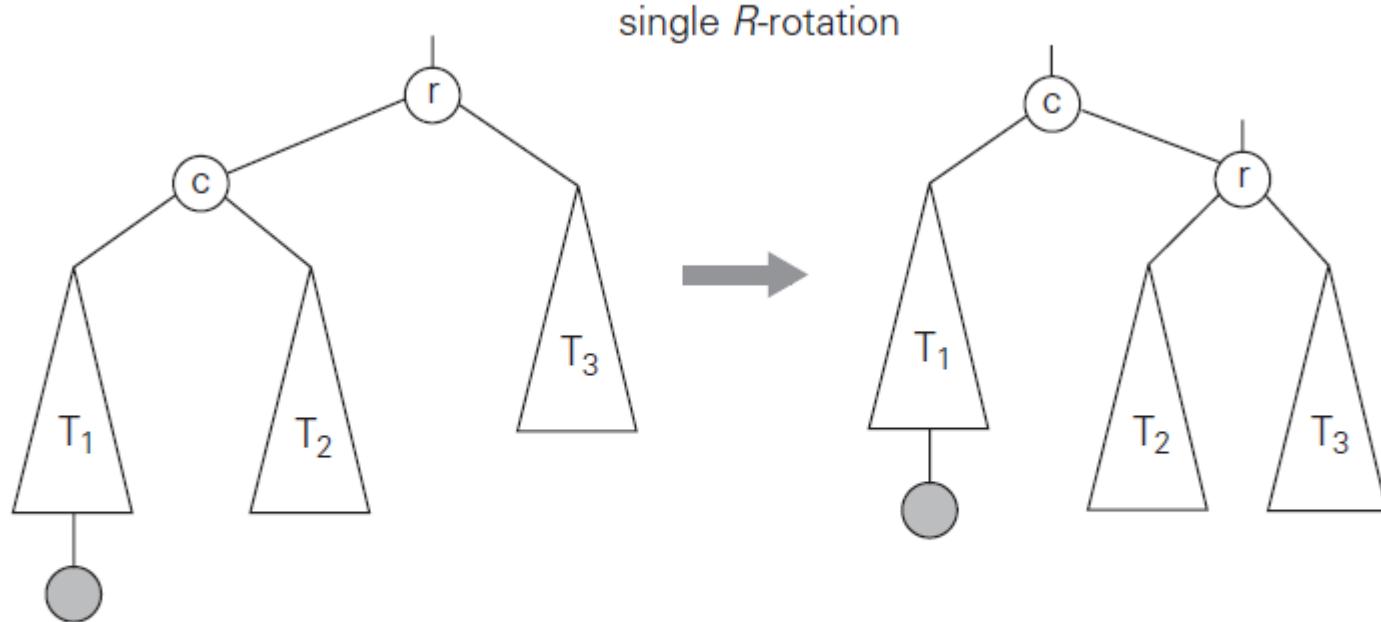


FIGURE 6.4 General form of the *R*-rotation in the AVL tree. A shaded node is the last one inserted.

General case: Double LR-rotation

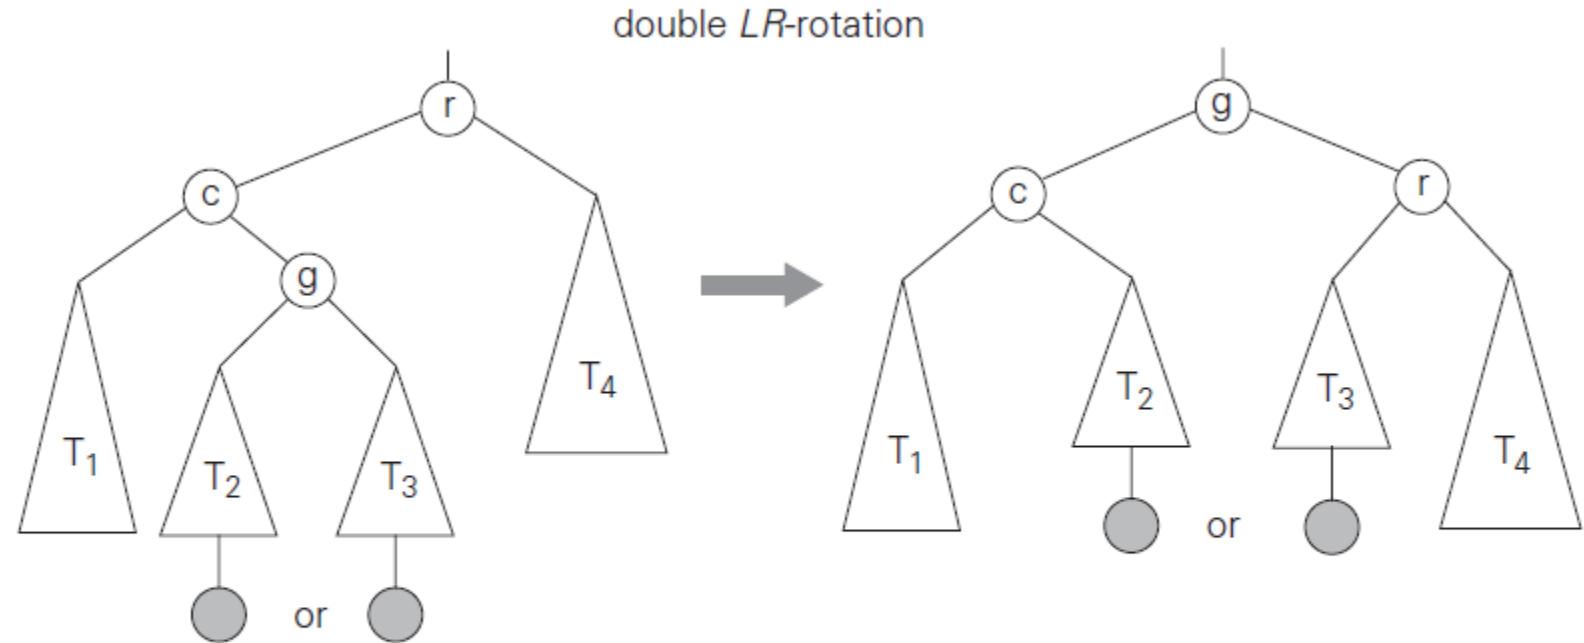
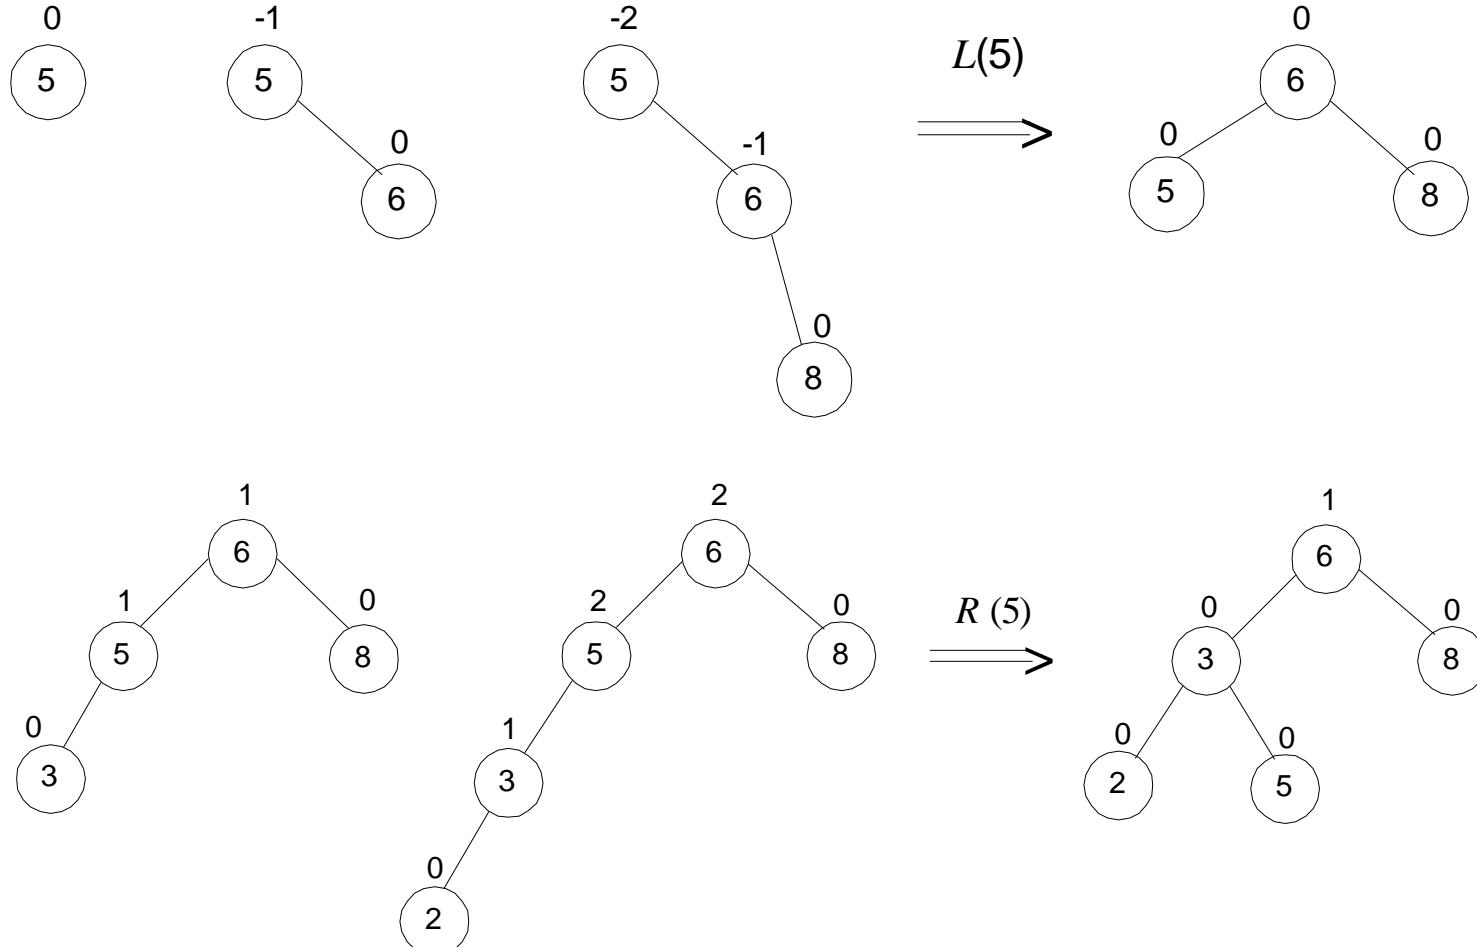


FIGURE 6.5 General form of the double *LR*-rotation in the AVL tree. A shaded node is the last one inserted. It can be either in the left subtree or in the right subtree of the root's grandchild.

AVL tree construction - an example



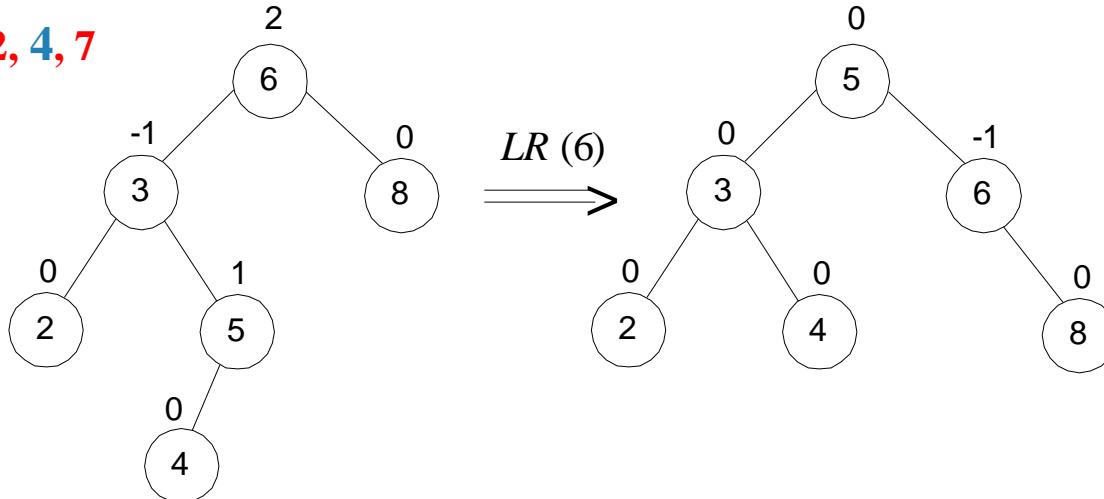
Construct an AVL tree for the list $5, 6, 8, 3, 2, 4, 7$



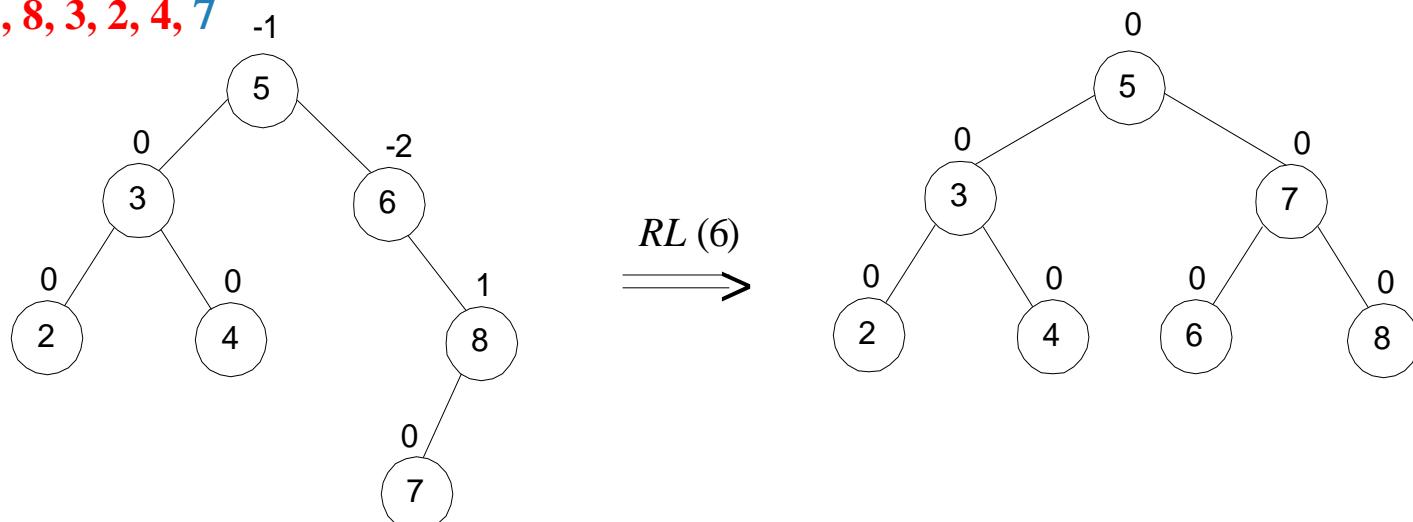
AVL tree construction - an example (cont.)



list: 5, 6, 8, 3, 2, 4, 7



list: 5, 6, 8, 3, 2, 4, 7



Analysis of AVL trees



ꝝ $h \leq 1.4404 \log_2(n + 2) - 1.3277$

average height: $1.01 \log_2 n + 0.1$ for large n (found empirically)

ꝝ Search and insertion are $O(\log n)$

ꝝ Deletion is more complicated but is also $O(\log n)$

ꝝ Disadvantages:

- frequent rotations
- complexity

ꝝ A similar idea: *red-black trees* (height of subtrees is allowed to differ by up to a factor of 2)

A. Levitin, "Introduction to the Design & Analysis of Algorithms," 3rd ed., Ch. 6 ©2012 Pearson

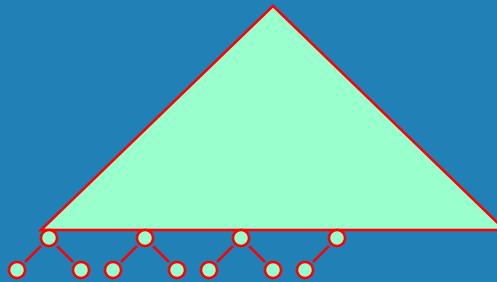
Education, Inc. Upper Saddle River, NJ. All Rights Reserved.

Heaps and Heapsort



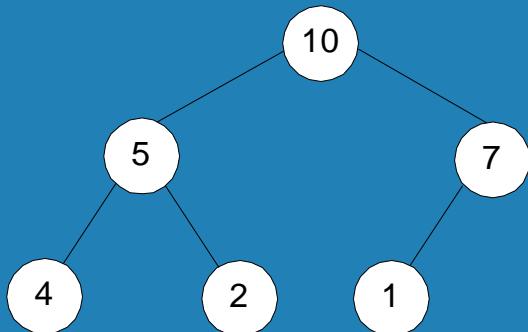
Definition A *heap* is a binary tree with keys at its nodes (one key per node) such that:

- ❑ It is essentially complete, i.e., all its levels are full except possibly the last level, where only some rightmost keys may be missing

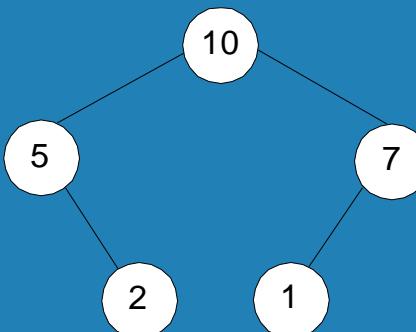


- ❑ The key at each node is \geq keys at its children

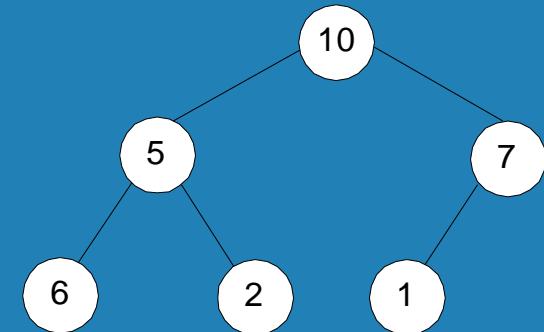
Illustration of the heap's definition



a heap



not a heap



not a heap

Note: Heap's elements are ordered top down (along any path down from its root), but they are not ordered left to right

Some Important Properties of a Heap



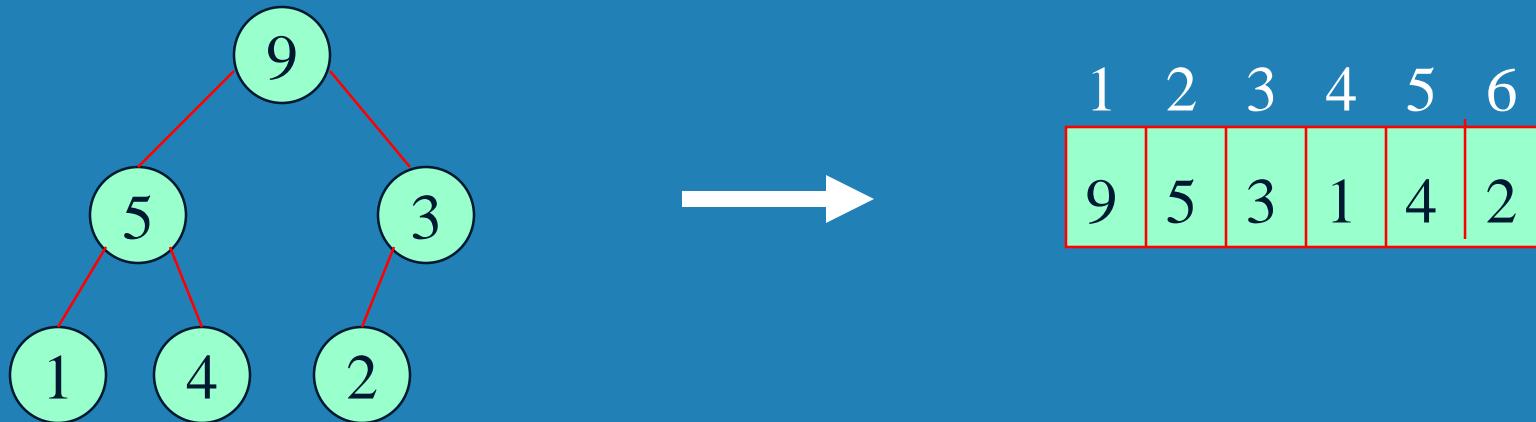
- ❑ Given n , there exists a unique binary tree with n nodes that is essentially complete, with $h = \lfloor \log_2 n \rfloor$
- ❑ The root contains the largest key
- ❑ The subtree rooted at any node of a heap is also a heap
- ❑ A heap can be represented as an array

Heap's Array Representation



Store heap's elements in an array (whose elements indexed, for convenience, 1 to n) in top-down left-to-right order

Example:



- ❑ Left child of node j is at $2j$
- ❑ Right child of node j is at $2j+1$
- ❑ Parent of node j is at $\lfloor j/2 \rfloor$
- ❑ Parental nodes are represented in the first $\lfloor n/2 \rfloor$ locations

Heap Construction (bottom-up)



Step 0: Initialize the structure with keys in the order given

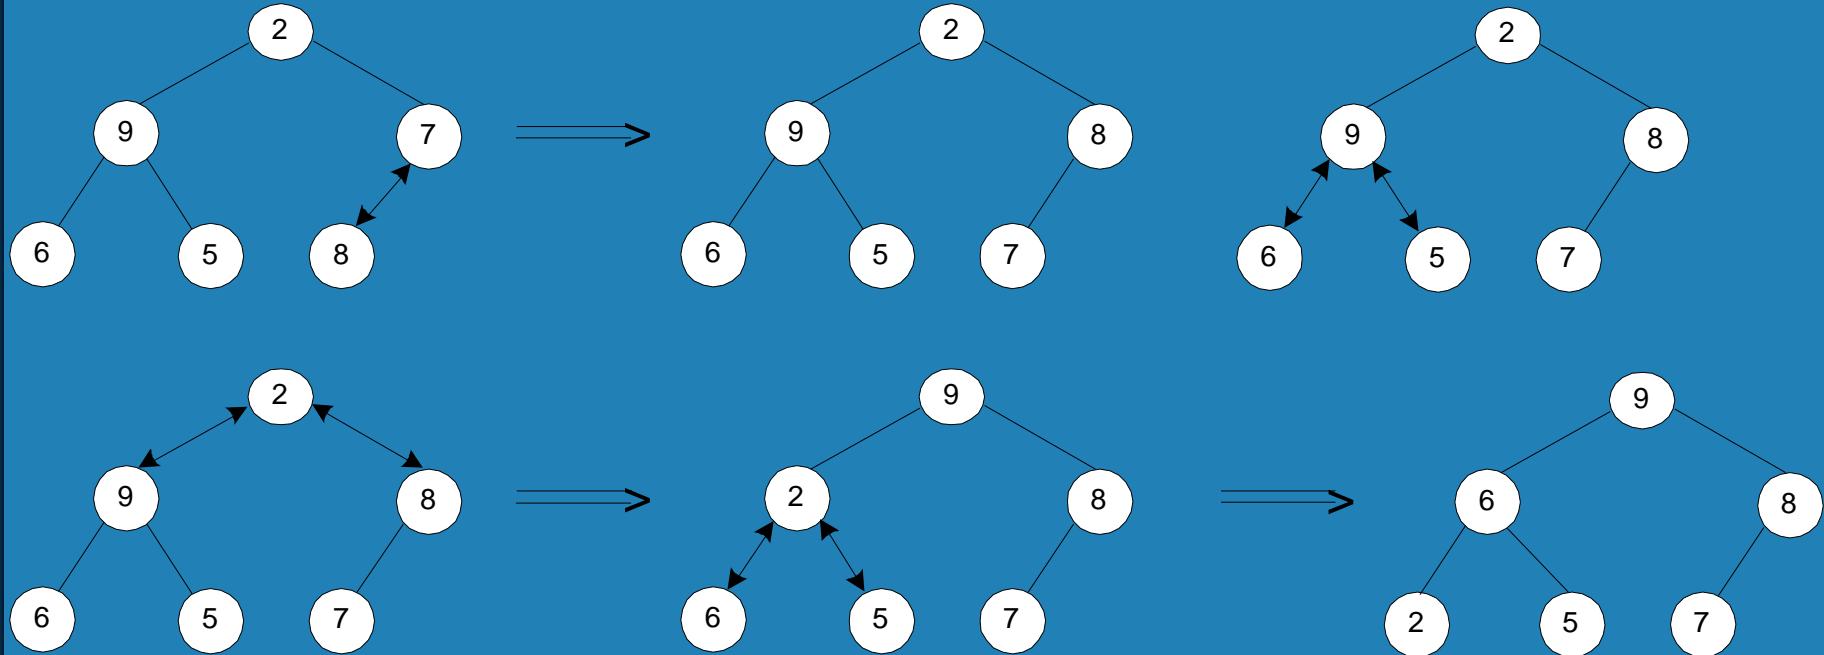
Step 1: Starting with the last (rightmost) parental node, fix the heap rooted at it, if it doesn't satisfy the heap condition: keep exchanging it with its largest child until the heap condition holds

Step 2: Repeat Step 1 for the preceding parental node

Example of Heap Construction



Construct a heap for the list 2, 9, 7, 6, 5, 8



Pseudopodia of bottom-up heap construction

```
Algorithm HeapBottomUp( $H[1..n]$ )
//Constructs a heap from the elements of a given array
// by the bottom-up algorithm
//Input: An array  $H[1..n]$  of orderable items
//Output: A heap  $H[1..n]$ 
for  $i \leftarrow \lfloor n/2 \rfloor$  downto 1 do
     $k \leftarrow i$ ;  $v \leftarrow H[k]$ 
     $heap \leftarrow \text{false}$ 
    while not  $heap$  and  $2 * k \leq n$  do
         $j \leftarrow 2 * k$ 
        if  $j < n$  //there are two children
            if  $H[j] < H[j + 1]$   $j \leftarrow j + 1$ 
        if  $v \geq H[j]$ 
             $heap \leftarrow \text{true}$ 
        else  $H[k] \leftarrow H[j]$ ;  $k \leftarrow j$ 
     $H[k] \leftarrow v$ 
```

Heapsort



Stage 1: Construct a heap for a given list of n keys

Stage 2: Repeat operation of root removal $n-1$ times:

- Exchange keys in the root and in the last (rightmost) leaf**
- Decrease heap size by 1**
- If necessary, swap new root with larger child until the heap condition holds**

Example of Sorting by Heapsort



Sort the list 2, 9, 7, 6, 5, 8 by heapsort

Stage 1 (heap construction)

1	9	<u>7</u>	6	5	8
2	<u>9</u>	8	6	5	7
<u>2</u>	9	8	6	5	7
9	<u>2</u>	8	6	5	7
9	6	8	2	5	7

Stage 2 (root/max removal)

<u>9</u>	6	8	2	5	7
7	6	8	2	5	9
<u>8</u>	6	7	2	5	9
5	6	7	2	8	9
<u>7</u>	6	5	2	8	9
2	6	5	7	8	9
<u>6</u>	2	5	7	8	9
5	2	6	7	8	9
<u>5</u>	2	6	7	8	9
2	5	6	7	8	9

Analysis of Heapsort



Stage 1: Build heap for a given list of n keys

worst-case

$$C(n) = \sum_{i=0}^{h-1} 2(h-i) 2^i = 2(n - \log_2(n+1)) \in \Theta(n)$$

/ # nodes at
 level i

Stage 2: Repeat operation of root removal $n-1$ times (fix heap)

worst-case

$$C(n) = \sum_{i=1}^{n-1} 2\log_2 i \in \Theta(n \log n)$$

Both worst-case and average-case efficiency: $\Theta(n \log n)$

In-place: yes

Stability: no (e.g., 1 1)

Priority Queue



A *priority queue* is the ADT of a set of elements with numerical priorities with the following operations:

- find element with highest priority
- delete element with highest priority
- insert element with assigned priority (see below)

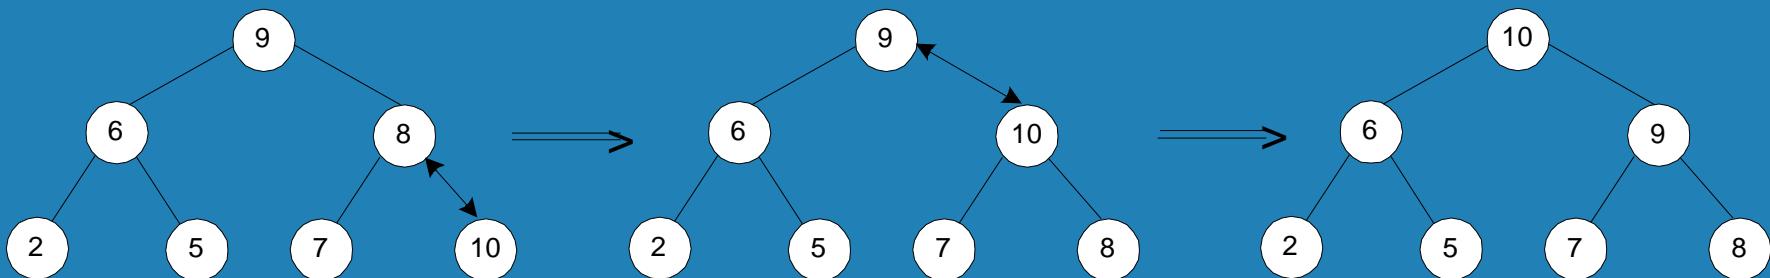
- ❑ Heap is a very efficient way for implementing priority queues
- ❑ Two ways to handle priority queue in which highest priority = smallest number

Insertion of a New Element into a Heap



- ❑ Insert the new element at last position in heap.
- ❑ Compare it with its parent and, if it violates heap condition, exchange them
- ❑ Continue comparing the new element with nodes up the tree until the heap condition is satisfied

Example: Insert key 10



Efficiency: $O(\log n)$